Collaborative Data Science Practices

Will Beasley

2019-02-01

Contents

Chapter 1

Prerequisites

This is a *sample* book written in **Markdown**. You can use anything that Pandoc's Markdown supports, e.g., a math equation $a^2 + b^2 = c^2$.

The **bookdown** package can be installed from CRAN or Github:

```r
install.packages("bookdown")
# or the development version
# devtools::install_github("rstudio/bookdown")
```

Remember each Rmd file contains one and only one chapter, and a chapter is defined by the first-level heading `#`.

To compile this example to PDF, you need XeLaTeX. You are recommended to install TinyTeX (which includes XeLaTeX): https://yihui.name/tinytex/.

Chapter 2

Architecture Principles

2.1 Encapsulation

2.2 Leverage team member's strenghts & avoid weaknesses

1. Focused code files
2. Metadata for content experts

2.3 Scales

1. Single source & single analysis
2. Multiple sources & multiple analyses

2.4 Consistency

1. Across Files {#consistency-files}
2. Across Languages
3. Across Projects

Chapter 3

Prototypical File

As stated before, in Consistency Files, using a consistent file structure can (a) improve the quality of the code because the structure has been proven over time to facilitate good practices and (b) allow your intentions to be more clear to teammates because they are familiar with the order and intentions of the chunks.

We use the term "chunk" for a section of code because it corresponds with knitr terminology (Xie, 2015), and in many cases, the chunk of our R file connects to a knitr Rmd file.

3.1   Clear Memory

3.2   Load Sources

3.3   Load Packages

3.4   Declare Globals

3.5   Load Data

3.6   Tweak Data

3.7   (Unique Content)

3.8   Verify Values

3.9   Specify Output Columns

3.10   Save to Disk or Database

Chapter 4

Prototypical Repository

https://github.com/wibeasley/RAnalysisSkeleton

4.1   Analysis

4.2   Data Public

1. Raw
2. Derived
3. Metadata
4. Database
5. Original

4.3   Data Unshared

4.4   Documentation

4.5   Manipulation

4.6   Stitched Output

4.7   Utility

Chapter 5

Data at Rest

5.1 Data States

1. Raw
2. Derived

    1. Project-wide File on Repo
    2. Project-wide File on Protected File Server
    3. User-specific File on Protected File Server
    4. Project-wide Database

3. Original

5.2 Data Containers

1. csv
2. rds
3. SQLite
4. Central Enterprise database
5. Central REDCap database
6. Containers to avoid for raw/input

    1. Proprietary like xlsx, sas7bdat

Chapter 6

Patterns

6.1 Ellis

6.1.1 Purpose

To incorporate outside data source into your system safely.

6.1.2 Philosophy

- Without data immigration, all warehouses are useless. Embrace the power of fresh information in a way that is:

  - repeatable when the datasource is updated (and you have to refresh your warehouse)
  - similar to other Ellis lanes (that are designed for other data sources) so you don't have to learn/remember an entirely new pattern. (Like Rubiks cube instructions.)

6.1.3 Guidelines

- Take small bites.

  - Like all software development, don't tackle all the complexity the first time. Start by processing only the important columns before incorporating move.
  - Use only the variables you need in the short-term, especially for new projects. As everyone knows, the variables from the upstream source can change. Don't spend effort writing code for variables you won't need for a few months/years; they'll likely change before you need them.
  - After a row passes through the `verify-values` chunk, you're accountable for any failures it causes in your warehouse. All analysts know that external data is messy, so don't be surprised. Sometimes I'll spend an hour writing an Ellis for 6 columns.

- Narrowly define each Ellis lane. One code file should strive to (a) consume only one CSV and (b) produce only one table. Exceptions include:

  1. if multiple input files are related, and really belong together (*e.g.*, one CSV per month, or one CSV per clinic). This scenario is pretty common.
  2. if the CSV should legitimately produce two different tables after munging. This happens infrequently, such as one warehouse table needs to be wide, and another long.

6.1.4   Examples

- https://github.com/wibeasley/RAnalysisSkeleton/blob/master/manipulation/te-ellis.R
- https://github.com/wibeasley/RAnalysisSkeleton/blob/master/manipulation/
- https://github.com/OuhscBbmc/usnavy-billets/blob/master/manipulation/survey-ellis.R

6.1.5   Elements

1. **Clear memory** In scripting languages like R (unlike compiled languages like Java), it's easy for old
   variables to hang around. Explicitly clear them before you run the file again.

   ```r
   rm(list=ls(all=TRUE)) #Clear the memory of variables from previous run. This is not called by knitr
   ```

2. **Load Sources** In R, a `source()`d file is run to execute its code. We prefer that a sourced file only
   load variables (like function definitions), instead of do real operations like read a dataset or perform a
   calculation. There are many times that you want a function to be available to multiple files in a repo;
   there are two approaches we like. The first is collecting those common functions into a single file (and
   then sourcing it in the callers). The second is to make the repo a legitimate R package.

   The first approach is better suited for quick & easy development. The second allows you to add
   documention and unit tests.

   ```r
   # ---- load-sources ------------------------------------------------------------
   source("./manipulation/osdh/ellis/common-ellis.R")
   ```

3. **Load Packages** This is another precaution necessary in a scripting language. Determine if the neces-
   sary packages are available on the machine. Avoiding attaching packages (with the `library()` function)
   when possible. Their functions don't need to be qualified (*e.g.*, `dplyr::intersect()`) and could cause
   naming conflicts. Even if you can guarantee they don't conflict with packages now, packages could
   add new functions in the future that do conflict.

   ```r
   # ---- load-packages -----------------------------------------------------------
   # Attach these package(s) so their functions don't need to be qualified: http://r-pkgs.had.co.nz/na
   library(magrittr            , quietly=TRUE)
   library(DBI                 , quietly=TRUE)

   # Verify these packages are available on the machine, but their functions need to be qualified: htt
   requireNamespace("readr"        )
   requireNamespace("tidyr"        )
   requireNamespace("dplyr"        ) # Avoid attaching dplyr, b/c its function names conflict with a l
   requireNamespace("testit")
   requireNamespace("checkmate")
   requireNamespace("OuhscMunge") #devtools::install_github(repo="OuhscBbmc/OuhscMunge")
   ```

4. **Declare Global Variables and Functions**. This includes defining the expected column names and
   types of the data sources; use `readr::cols_only()` (as opposed to `readr::cols()`) to ignore any new
   columns that may be been added since the dataset's last refresh.

   ```r
   # ---- declare-globals ---------------------------------------------------------
   ```

5. **Load Data Source(s)** Read all data (*e.g.*, database table, networked CSV, local lookup table). After
   this chunk, no new data should be introduced. This is for the sake of reducing human cognition load.
   Everything below this chunk is derived from these first four chunks.

```
# ---- load-data -----------------------------------------------------------------
```

6. **Tweak Data**

```
# ---- tweak-data ----------------------------------------------------------------
```

7. **Body of the Ellis**

8. **Verify**

```
# ---- verify-values -------------------------------------------------------------
county_month_combo   <- paste(ds$county_id, ds$month)
checkmate::assert_character(county_month_combo, pattern ="^\\d{1,2} \\d{4}-\\d{2}-\\d{2}$", any.mi
```

9. **Specify Columns** Define the exact columns and order to upload to the database. Once you import a column into a warehouse that multiple people are using, it's tough to remove it.

```
# ---- specify-columns-to-upload -------------------------------------------------
```

10. **Welcome** into your warehouse. Until this chunk, nothing should be persisted.

```
# ---- upload-to-db --------------------------------------------------------------
# ---- save-to-disk --------------------------------------------------------------
```

6.2   Arch

6.3   Ferry

6.4   Scribe

6.5   Analysis

6.6   Presentation -Static

6.7   Presentation -Interactive

6.8   Metadata

Chapter 7

Security & Private Data

7.1 File-level permissions

7.2 Database permissions

7.3 Public & Private Repositories

7.3.1 Scrubbing GitHub history

Occassionaly files may be committed to your git repository that need to be removed completely. Not just from the current collections of files (*i.e.*, the branch's head), but from the entire history of the repo.

Scrubbing is require typically when (a) a sensitive file has been accidentally commited and pushed to GitHub, or (b) a huge file has bloated your repository and disrupted productivity.

The two suitable scrubbing approaches both require the command line. The first is the `git-filter-branch` command within git, and the second is the BFG repo-cleaner. We use the second approach, which is [recommended by GitHub]; it requires 15 minutes to install and configure from scratch, but then is much easier to develop against, and executes much faster.

The bash-centric steps below remove any file*s* from the repo history called 'monster-data.csv' from the 'bloated' repository.

1. If the file contains passwords, change them immediately.

2. Delete 'monster-data.csv' from your branch and push the commit to GitHub.

3. Ask your collaborators to push any outstanding commits to GitHub and delete their local copy of the repo. Once scrubbing is complete, they will re-clone it.

4. Download and install the most recent Java JRE from the Oracle site.

5. Download the most recent jar file from the BFG site to the home directory.

6. Clone a fresh copy of the repository in the user's home directory. The `--mirror` argument avoids downloading every file, and downloads only the bookkeeping details required for scrubbing.

```
cd ~
git clone --mirror https://github.com/your-org/bloated.git
```

7. Remove all files (in any directory) called 'monster-data.csv'.

```
java -jar bfg-*.jar --delete-files monster-data.csv bloated.git
```

8. Reflog and garbage collect the repo.

```
cd bloated.git
git reflog expire --expire=now --all && git gc --prune=now --aggressive
```

9. Push your local changes to the GitHub server.

```
git push
```

10. Delete the bfg jar from the home directory.

```
cd ~
rm bfg-*.jar
```

11. Ask your collaborators to reclone the repo to their local machine. It is important they restart with a fresh copy, so the once-scrubbed file is not reintroduced into the repo's history.

12. If the file contains sensitive information, like passwords or PHI, ask GitHub to refresh the cache so the file's history isn't accessible through their website, even if the repo is private.

7.3.1.0.1   Resources

- BFG Repo-Cleaner site
- Additional BFG instructions
- GitHub Sensitive Data Removal Policy

Chapter 8

Automation

Chapter 9

Scaling Up

9.1   Data Storage

1. Local File vs Conventional Database vs Redshift
2. Usage Cases

9.2   Data Processing

1. R vs SQL
2. R vs Spark

Chapter 10

Parallel Collaboration

10.1 Social Contract

1. Issues
2. Organized Commits & Coherent Diffs
3. Branch & Merge Strategy

10.2 Code Reviews

1. Daily Reviews of PRs
2. Periodic Reviews of Files

10.3 Remote

1. Headset & sharing screens

Chapter 11

Documentation

11.1 Team-wide

11.2 Project-specific

11.3 Dataset Origin & Structure

11.4 Issues & Tasks

11.5 Flow Diagrams

11.6 Setting up new machine

(example)

Chapter 12

Publishing Results

12.1   To Other Analysts

12.2   To Researchers & Content Experts

12.3   To Technical-Phobic Audiences

Chapter 13

Testing, Validation, & Defensive Programming

13.1 Testing Functions

13.2 Defensive Programming

1. Throwing errors

13.3 Validator

1. Benefits for Analysts
2. Benefits for Data Collectors

Chapter 14

Troubleshooting and Debugging

14.1 Finding Help

1. Within your group (eg, Thomas and REDCap questions)
2. Within your university (eg, SCUG)
3. Outside (eg, Stack Overflow; GitHub issues)

14.2 Debugging

1. `traceback()`, `browser()`, etc

Chapter 15

Considerations when Selecting Tools

https://github.com/OuhscBbmc/RedcapExamplesAndPatterns/blob/master/DocumentationGlobal/ResourcesInstallation.md

15.1 Required Installation

15.2 Recommended Installation

15.3 Optional Installation

15.4 Asset Locations

Chapter 16

Considerations when Selecting Tools

16.1 General

16.1.1 The Component's Goal

While disussing the advantages and disadvanages of tools, a colleague once said, "Tidyverse packages don't do anything that I can't already do in Base R, and sometimes it even requires more lines of code". Regardless if I agree, I feel these two points are irrelevant. Sometimes the advantage of a tool isn't to expand existing capabilities, but rather to facilitate development and maintaince for the same capability.

Likewise, I care less about the line count, and more about the readability. I'd prefer to maintain a 20-line chunk that is familar and readable than a 10-line chunk with dense phrases and unfamiliar functions. The bottleneck for most of our projects is human time, not execution time.

16.1.2 Current Skillset of Team

16.1.3 Desired Future Skillset of Team

16.1.4 Skillset of Audience

16.2 Languages

16.3 R Packages

16.4 Database

Chapter 17

Growing a Team

17.1 Recruiting

17.2 Training to Data Science

1. Starting with a Researcher
2. Starting with a Statistician
3. Starting with a DBA
4. Starting with a Software Developer

17.3 Bridges Outside the Team

1. Monthly User Groups
2. Annual Conferences

Chapter 18

Introduction

You can label chapter and section titles using `{#label}` after them, e.g., we can reference Chapter 18. If you do not manually label them, there will be automatic labels anyway, e.g., Chapter 2.

Figures and tables with captions will be placed in `figure` and `table` environments, respectively.

```r
par(mar = c(4, 4, .1, .1))
plot(pressure, type = 'b', pch = 19)
```

Reference a figure by its code chunk label with the `fig:` prefix, e.g., see Figure 18.1. Similarly, you can reference tables generated from `knitr::kable()`, e.g., see Table 18.1.

```r
knitr::kable(
  head(iris, 20), caption = 'Here is a nice table!',
  booktabs = TRUE
)
```

You can write citations, too. For example, we are using the **bookdown** package (Xie, 2018) in this sample book, which was built on top of R Markdown and **knitr** (Xie, 2015).
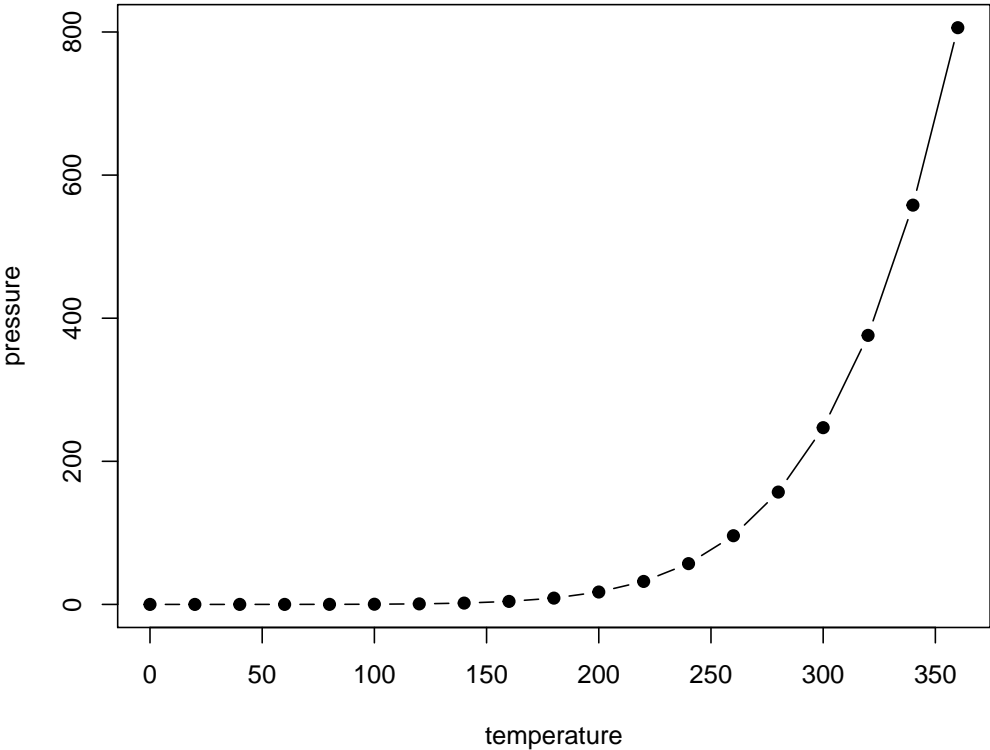
Figure 18.1: Here is a nice figure!

Table 18.1: Here is a nice table!

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
5.1	3.5	1.4	0.2	setosa
4.9	3.0	1.4	0.2	setosa
4.7	3.2	1.3	0.2	setosa
4.6	3.1	1.5	0.2	setosa
5.0	3.6	1.4	0.2	setosa
5.4	3.9	1.7	0.4	setosa
4.6	3.4	1.4	0.3	setosa
5.0	3.4	1.5	0.2	setosa
4.4	2.9	1.4	0.2	setosa
4.9	3.1	1.5	0.1	setosa
5.4	3.7	1.5	0.2	setosa
4.8	3.4	1.6	0.2	setosa
4.8	3.0	1.4	0.1	setosa
4.3	3.0	1.1	0.1	setosa
5.8	4.0	1.2	0.2	setosa
5.7	4.4	1.5	0.4	setosa
5.4	3.9	1.3	0.4	setosa
5.1	3.5	1.4	0.3	setosa
5.7	3.8	1.7	0.3	setosa
5.1	3.8	1.5	0.3	setosa

Chapter 19

Scratch Pad of Loose Ideas

19.1   Chapters & Sections to Form

1. Tools to Consider
    1. tidyverse
    2. odbc

Bibliography

Xie, Y. (2015). *Dynamic Documents with R and knitr.* Chapman and Hall/CRC, Boca Raton, Florida, 2nd edition. ISBN 978-1498716963.

Xie, Y. (2018). *bookdown: Authoring Books and Technical Documents with R Markdown.* R package version 0.9.